

# HPML Assignment 2

## Program Outputs

### Program C1

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import argparse
import time

class BasicBlock(nn.Module):
    def __init__(self, in_channels, out_channels, stride=1):
        super(BasicBlock, self).__init__()
        self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=3, stride=stride, padding=1, bias=False)
        self.bn1 = nn.BatchNorm2d(out_channels)
        self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size=3, stride=1, padding=1, bias=False)
        self.bn2 = nn.BatchNorm2d(out_channels)

        self.shortcut = nn.Sequential()
        if stride != 1 or in_channels != out_channels:
            self.shortcut = nn.Sequential(
                nn.Conv2d(in_channels, out_channels, kernel_size=1, stride=stride, bias=False),
                nn.BatchNorm2d(out_channels)
            )

    def forward(self, x):
        out = F.relu(self.bn1(self.conv1(x)))
        out = self.bn2(self.conv2(out))
        out += self.shortcut(x)
        out = F.relu(out)
        return out

class ResNet18(nn.Module):
    def __init__(self):
        super(ResNet18, self).__init__()
        self.conv1 = nn.Conv2d(3, 64, kernel_size=3, stride=1, padding=1, bias=False)
        self.bn1 = nn.BatchNorm2d(64)

        self.layer1 = self._make_layer(64, 64, 2, stride=1)
        self.layer2 = self._make_layer(64, 128, 2, stride=2)
```

```

        self.layer3 = self._make_layer(128, 256, 2, stride=2)
        self.layer4 = self._make_layer(256, 512, 2, stride=2)

        self.linear = nn.Linear(512, 10)

    def _make_layer(self, in_channels, out_channels, blocks, stride):
        layers = []
        layers.append(BasicBlock(in_channels, out_channels, stride))
        for _ in range(1, blocks):
            layers.append(BasicBlock(out_channels, out_channels))
        return nn.Sequential(*layers)

    def forward(self, x):
        out = F.relu(self.bn1(self.conv1(x)))
        out = self.layer1(out)
        out = self.layer2(out)
        out = self.layer3(out)
        out = self.layer4(out)
        out = F.avg_pool2d(out, 4)
        out = out.view(out.size(0), -1)
        out = self.linear(out)
        return out

def main():
    parser = argparse.ArgumentParser(description='Train ResNet-18 on CIFAR-10')
    parser.add_argument('--use_cuda', action='store_true', help='Use CUDA if available')
    parser.add_argument('--data_path', type=str, default='./data', help='Path to CIFAR-10 data')
    parser.add_argument('--num_workers', type=int, default=2, help='Number of data loader workers')
    parser.add_argument('--optimizer', type=str, default='sgd', help='Optimizer to use (sgd, adam)')
    args = parser.parse_args()

    device = torch.device('cuda' if args.use_cuda and torch.cuda.is_available() else 'cpu')

    # DataLoader for CIFAR-10 dataset
    transform = transforms.Compose([
        transforms.RandomCrop(32, padding=4),
        transforms.RandomHorizontalFlip(p=0.5),
        transforms.ToTensor(),
        transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010))
    ])

    train_dataset = torchvision.datasets.CIFAR10(root=args.data_path, train=True, download=True)
    train_loader = DataLoader(train_dataset, batch_size=128, shuffle=True, num_workers=args.num_workers)

    # Model, loss, and optimizer
    model = ResNet18().to(device)

```

```

criterion = nn.CrossEntropyLoss()

if args.optimizer.lower() == 'sgd':
    optimizer = optim.SGD(model.parameters(), lr=0.1, momentum=0.9, weight_decay=5e-4)
elif args.optimizer.lower() == 'sgd_nesterov':
    optimizer = optim.SGD(model.parameters(), lr=0.1, momentum=0.9, weight_decay=5e-4, nesterov=True)
elif args.optimizer.lower() == 'adagrad':
    optimizer = optim.Adagrad(model.parameters(), lr=0.01, weight_decay=5e-4)
elif args.optimizer.lower() == 'adadelat':
    optimizer = optim.Adadelta(model.parameters(), lr=1.0, weight_decay=5e-4)
elif args.optimizer.lower() == 'adam':
    optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=5e-4)
else:
    raise ValueError('Unsupported optimizer. Use "sgd", "sgd_nesterov", "adagrad", "adadelat", "adam"')

# Training loop
num_epochs = 5
for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0
    correct = 0
    total = 0

    data_loading_time = 0.0
    training_time = 0.0

    # Start timing for data loading
    if device.type == 'cuda':
        torch.cuda.synchronize()
    data_loading_start_time = time.perf_counter()

    for batch_idx, (inputs, targets) in enumerate(train_loader):
        if device.type == 'cuda':
            torch.cuda.synchronize()
        data_loading_end_time = time.perf_counter()
        data_loading_time += data_loading_end_time - data_loading_start_time

        # Start timing for training
        if device.type == 'cuda':
            torch.cuda.synchronize()
        training_start_time = time.perf_counter()

        inputs, targets = inputs.to(device), targets.to(device)
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, targets)

```

```

        loss.backward()
        optimizer.step()

    if device.type == 'cuda':
        torch.cuda.synchronize()
        training_end_time = time.perf_counter()
        training_time += training_end_time - training_start_time

    running_loss += loss.item()
    _, predicted = outputs.max(1)
    total += targets.size(0)
    correct += predicted.eq(targets).sum().item()

    if batch_idx % 100 == 0:
        print(f'Epoch [{epoch+1}/{num_epochs}], Batch [{batch_idx+1}/{len(train_loader)}]')

        # Reset data loading timer
        if device.type == 'cuda':
            torch.cuda.synchronize()
            data_loading_start_time = time.perf_counter()

        # Total epoch time
        if device.type == 'cuda':
            torch.cuda.synchronize()
            epoch_end_time = time.perf_counter()
            total_epoch_time = data_loading_time + training_time

        print(f'Epoch [{epoch+1}/{num_epochs}] Summary: Data Loading Time: {data_loading_time}')

if __name__ == '__main__':
    main()

```

## Program C2

```
bash-5.1$ python c2.py --num_workers 2 --optimizer sgd
```

Files already downloaded and verified

```

Epoch [1/5], Batch [1/391], Loss: 2.4453, Accuracy: 7.81%
Epoch [1/5], Batch [101/391], Loss: 2.3240, Accuracy: 23.24%
Epoch [1/5], Batch [201/391], Loss: 2.0366, Accuracy: 28.61%
Epoch [1/5], Batch [301/391], Loss: 1.9009, Accuracy: 31.97%
Epoch [1/5] Summary: Data Loading Time: 0.7442s, Training Time: 201.4467s, Total Epoch Time: 202.1909s
Epoch [2/5], Batch [1/391], Loss: 1.3482, Accuracy: 47.66%
Epoch [2/5], Batch [101/391], Loss: 1.4344, Accuracy: 47.25%
Epoch [2/5], Batch [201/391], Loss: 1.3928, Accuracy: 48.81%
Epoch [2/5], Batch [301/391], Loss: 1.3475, Accuracy: 50.94%

```

```

Epoch [2/5] Summary: Data Loading Time: 0.7746s, Training Time: 184.8464s, Total Epoch Time:
Epoch [3/5], Batch [1/391], Loss: 1.0872, Accuracy: 64.06%
Epoch [3/5], Batch [101/391], Loss: 1.0842, Accuracy: 60.99%
Epoch [3/5], Batch [201/391], Loss: 1.0535, Accuracy: 62.57%
Epoch [3/5], Batch [301/391], Loss: 1.0470, Accuracy: 62.92%
Epoch [3/5] Summary: Data Loading Time: 0.8489s, Training Time: 194.0466s, Total Epoch Time:
Epoch [4/5], Batch [1/391], Loss: 0.7287, Accuracy: 72.66%
Epoch [4/5], Batch [101/391], Loss: 0.9070, Accuracy: 67.71%
Epoch [4/5], Batch [201/391], Loss: 0.8906, Accuracy: 68.35%
Epoch [4/5], Batch [301/391], Loss: 0.8722, Accuracy: 69.12%
Epoch [4/5] Summary: Data Loading Time: 0.8840s, Training Time: 190.8991s, Total Epoch Time:
Epoch [5/5], Batch [1/391], Loss: 0.8141, Accuracy: 73.44%
Epoch [5/5], Batch [101/391], Loss: 0.7756, Accuracy: 72.79%
Epoch [5/5], Batch [201/391], Loss: 0.7560, Accuracy: 73.34%
Epoch [5/5], Batch [301/391], Loss: 0.7423, Accuracy: 73.95%
Epoch [5/5] Summary: Data Loading Time: 0.8239s, Training Time: 196.3094s, Total Epoch Time:

```

### Program C3

```
bash-5.1$ python c3.py
```

```

Running with 0 workers...
Data loading time for 0 workers: 10.915s
Running with 4 workers...
Data loading time for 4 workers: 0.7999s
Running with 8 workers...
Data loading time for 8 workers: 0.8983s

```

Summary of Results:

```

Number of Workers: 0, Data Loading Time: 10.91s
Number of Workers: 4, Data Loading Time: 0.80s
Number of Workers: 8, Data Loading Time: 0.90s

```

The best number of workers for runtime performance is: 4

Graph: The graph basically shows that there is diminishing returns with increase in number of workers.

### Program C4

```
bash-5.1$ python c4.py
```

```

Running with 1 worker...
Running with 4 workers...
1 Worker - Data Loading Time: 0.70s, Training Time: 177.81s
4 Workers - Data Loading Time: 0.78s, Training Time: 189.10s

```

The data loading times are similar or even slightly decreasing when using 4 workers compared to 1 worker. The benefits of multiple workers might not be noticeable with such a small dataset.

### Program C5

```
bash-5.1$ python c5.py
```

```
Running training with GPU (using 4 workers)...
```

```
Running training with CPU (using 4 workers)...
```

Results over 5 epochs:

Average GPU training time per epoch: 5.31 seconds

Average CPU training time per epoch: 385.19 seconds

Overall speedup factor: 72.51x

Epoch-wise comparison:

Epoch	GPU Time (s)	CPU Time (s)	Speedup
-------	--------------	--------------	---------

1	7.80	379.64	48.69x
2	4.64	382.03	82.29x
3	4.70	398.94	84.93x
4	4.79	382.11	79.77x
5	4.63	383.23	82.70x

### Program C6

```
bash-5.1$ python c6.py
```

```
Running training with SGD...
```

```
Running training with SGD_NESTEROV...
```

```
Running training with ADAGRAD...
```

```
Running training with ADADELTA...
```

```
Running training with ADAM...
```

Summary (averaged over 5 epochs):

Optimizer	Training Time	Loss	Accuracy
ADADELTA	15.2735	1.9484	27.778
ADAGRAD	14.9701	1.9013	29.016
ADAM	15.0567	1.7274	36.856
SGD	15.2735	2.3505	21.024
SGD_NESTEROV	14.8702	2.1111	27.012

Detailed Results:

SGD:

Epoch	Training Time (s)	Loss	Accuracy (%)
-------	-------------------	------	--------------

1	17.35	2.5339	10.94
1	14.74	2.7230	17.44
1	14.78	2.3524	21.67
1	14.74	2.1909	24.60
1	14.75	1.9521	30.47

SGD\_NESTEROV:

Epoch	Training Time (s)	Loss	Accuracy (%)
1	15.13	2.5006	8.59
1	14.82	2.5303	19.66
1	14.81	2.1639	25.98
1	14.79	1.9974	30.05
1	14.80	1.3634	50.78

ADAGRAD:

Epoch	Training Time (s)	Loss	Accuracy (%)
1	15.41	2.3631	8.59
1	14.85	2.0576	24.95
1	14.86	1.8760	30.85
1	14.85	1.7731	34.60
1	14.88	1.4369	46.09

ADADELTA:

Epoch	Training Time (s)	Loss	Accuracy (%)
1	15.63	2.3647	10.16
1	15.19	2.0376	25.40
1	15.15	1.8836	30.14
1	15.19	1.7792	34.13
1	15.20	1.6769	39.06

ADAM:

Epoch	Training Time (s)	Loss	Accuracy (%)
1	15.41	2.4420	8.59
1	14.93	1.8183	31.90
1	14.98	1.6571	38.13
1	14.97	1.5526	42.38
1	15.00	1.1672	63.28

## Program C7

```
bash-5.1$ python c7.py --use_cuda --num_workers 4 --optimizer sgd
Files already downloaded and verified
Epoch [1/5], Batch [1/391], Loss: 2.3068, Accuracy: 8.59%
Epoch [1/5], Batch [101/391], Loss: 2.2537, Accuracy: 14.02%
Epoch [1/5], Batch [201/391], Loss: 2.1501, Accuracy: 18.73%
Epoch [1/5], Batch [301/391], Loss: 2.0567, Accuracy: 22.46%
Epoch [1/5] Summary: Data Loading Time: 0.7678s, Training Time: 3.1920s, Total Epoch Time: 3.9598s
Epoch [2/5], Batch [1/391], Loss: 1.7660, Accuracy: 31.25%
Epoch [2/5], Batch [101/391], Loss: 1.6879, Accuracy: 37.23%
Epoch [2/5], Batch [201/391], Loss: 1.6605, Accuracy: 38.36%
Epoch [2/5], Batch [301/391], Loss: 1.6239, Accuracy: 39.80%
Epoch [2/5] Summary: Data Loading Time: 0.8277s, Training Time: 2.4362s, Total Epoch Time: 3.2639s
Epoch [3/5], Batch [1/391], Loss: 1.4339, Accuracy: 42.19%
Epoch [3/5], Batch [101/391], Loss: 1.4382, Accuracy: 47.80%
Epoch [3/5], Batch [201/391], Loss: 1.4222, Accuracy: 48.47%
Epoch [3/5], Batch [301/391], Loss: 1.3870, Accuracy: 49.65%
Epoch [3/5] Summary: Data Loading Time: 0.7970s, Training Time: 2.5720s, Total Epoch Time: 3.3690s
Epoch [4/5], Batch [1/391], Loss: 1.3065, Accuracy: 50.78%
Epoch [4/5], Batch [101/391], Loss: 1.2422, Accuracy: 55.32%
Epoch [4/5], Batch [201/391], Loss: 1.2311, Accuracy: 56.01%
Epoch [4/5], Batch [301/391], Loss: 1.1983, Accuracy: 57.30%
Epoch [4/5] Summary: Data Loading Time: 0.8031s, Training Time: 2.8857s, Total Epoch Time: 3.6888s
Epoch [5/5], Batch [1/391], Loss: 1.0690, Accuracy: 57.81%
Epoch [5/5], Batch [101/391], Loss: 1.0938, Accuracy: 61.68%
Epoch [5/5], Batch [201/391], Loss: 1.0752, Accuracy: 62.27%
Epoch [5/5], Batch [301/391], Loss: 1.0434, Accuracy: 63.37%
Epoch [5/5] Summary: Data Loading Time: 0.7971s, Training Time: 2.4515s, Total Epoch Time: 3.2516s
```

## Answers of Theroetical Questions

### Question 1

How many convolutional layers are in the ResNet-18 model?

### Answer 1

- 1 initial convolutional layer
- 8 Basic Blocks (2 in each of the 4 layers) \* 2 conv layers per block = 16 conv layers
- 3 shortcut convolutional layers (in the first block of layers 2, 3, and 4)
- Total of 20 convolutional layers

### Question 2

What is the input dimension of the last linear layer?



**Answer 2**

```
self.linear = nn.Linear(512, 10)
```

Input dimension is 512

**Question 3**

How many trainable parameters and how many gradients in the ResNet-18 model that you build (please show both the answer and the code that you use to count them), when using SGD optimizer?

**Answer 3**

- Total parameters: 11173962
- Trainable parameters: 11173962
- Total gradients: 11173962

Refer file q3.py for code

**Question 4**

How many trainable parameters and how many gradients in the ResNet-18 model that you build (please show only the answer), when using ADAM optimizer?

**Answer 4**

- Total parameters: 11173962
- Trainable parameters: 11173962
- Total gradients: 11173962

Refer file q4.py for code